

Computational statistics, lecture 4

Frank Miller, Department of Computer and Information Science,
Linköping University
November 21, 2023

Markov chain Monte Carlo (MCMC)

(Literature: Gentle, 7.3-7.4; Givens and Hoeting, 7.1, 7.3)

- The algorithms considered so far generate sequences of **independent** observations which follow the target distribution exactly
- We will now consider a method which generates a sequence of **dependent** observations which follow the target distribution **approximately**
- The next observation ($t+1$) will be generated based on a proposal distribution g which depends on the current observation (t), i.e. $g(\cdot | X^{(t)})$
- Since $X^{(t+1)}$ depends on $X^{(t)}$ but not on earlier observations, the sequence $(X^{(t)})$ is a Markov chain

MCMC – Metropolis-Hastings algorithm

- A general method to generate the Markov chain is the Metropolis-Hastings (MH) algorithm
- A starting value $x^{(0)}$ is generated from some starting distribution
- Given observation $x^{(t)}$, generate $x^{(t+1)}$ as follows:

1. Sample a candidate x^* from a proposal distribution $g(\cdot|x^{(t)})$

2. Compute the MH ratio $R(x^{(t)}, x^*) = \frac{f(x^*) g(x^{(t)}|x^*)}{f(x^{(t)}) g(x^*|x^{(t)})}$

3. Sample $x^{(t+1)}$ according to

$$x^{(t+1)} = \begin{cases} x^*, & \text{with probability } \min\{R(x^{(t)}, x^*), 1\} \\ x^{(t)}, & \text{otherwise} \end{cases}$$

4. If more observations needed, set $t \leftarrow t+1$; go to 1

Metropolis algorithm


Special case when g is symmetric:

$$g(x^*|x^{(t)}) = g(x^{(t)}|x^*)$$

$$= \frac{f(x^*)}{f(x^{(t)})}$$

Metropolis alg. – Ex.1

- For illustration, we consider two-dimensional distribution with density f according to contour lines in figure
- Proposal distribution
$$g(x^* | x^{(t)}) = g(x^{(t)} | x^*)$$
$$= \frac{1}{\pi r^2} \mathbf{1}\{\|x^{(t)} - x^*\| < r\}$$
for some constant r (here=1)




Metropolis alg. – Ex.1

- Proposal distribution


$$g(x^* | x^{(t)}) = g(x^{(t)} | x^*)$$

$$= \frac{1}{\pi r^2} \mathbf{1}\{\|x^{(t)} - x^*\| < r\}$$
 for some constant r (here=1)
- Start here with $x^{(0)} = (1, -0.5)$
- Randomize uniformly on unit circle around $x^{(0)}$ (proposal distribution); result $x^* = (0.58, 0.08)$
- $f(x^*) = 0.296 > f(x^{(0)}) = 0.098$; so this was an uphill step and is automatically accepted ($R(x^{(t)}, x^*) = \frac{f(x^*)}{f(x^{(t)})} > 1$)




Metropolis alg. – Ex.1

- $x^{(0)} = (1, -0.5)$
- Uphill steps: $x^{(1)} = (0.58, 0.08)$
- $x^{(2)} = (-0.33, 0.13)$
- $x^{(3)} = (-0.23, 0.05)$
- Then downhill step proposed:
 $x^* = (-0.32, 0.4)$,
 $R(x^{(t)}, x^*) = \frac{f(x^*)}{f(x^{(t)})} = 0.774$
- Random Unif(0,1) generated: 0.573 and since this is smaller than $R=0.774$, $x^{(4)} = x^* = (-0.32, 0.4)$ is accepted
- Again downhill step proposed: $x^* = (-0.67, 1.31)$, $R(x^{(t)}, x^*) = \frac{f(x^*)}{f(x^{(t)})} = 0.560$; random Unif(0,1): 0.890 and rejection of x^*
- $x^{(5)} = x^{(4)} = (-0.32, 0.4)$




Metropolis alg. – Ex.1

- After several additional iterations (see red lines for rejected proposals), one part of the distribution was explored to a good extend
- Since uphill steps preferred, part of distribution with local maximum at $(-0.5, -0.5)$ is not yet "detected" at all
- Occasionally, the path will arrive at this part as well



Metropolis alg. – Ex.1

- Now, larger parts of distribution explored



- A couple of animations can be found on:
<https://chi-feng.github.io/mcmc-demo/app.html#RandomWalkMH,standard>
(choose Algorithm: RandomWalkMH)

Convergence of Metropolis-Hastings

- If Metropolis-Hastings generated sequence $(X^{(t)})$ is an **irreducible and aperiodic chain** (compare Lecture LM2), the distribution of $(X^{(t)})$ **converges to target distribution**
- For example, if target distribution is uniform distribution on intervals $[0, 1/2]$ and $[3/2, 2]$, and proposal distribution is uniform distribution on $[X^{(t)}-1/2, X^{(t)}+1/2]$, the requirements above are violated

Bayesian analysis

- Data y is collected and assumed that it is generated according to a distribution with density $f(y|\theta)$; θ is a parameter(-vector) to be estimated
- The posterior density is proportional to product of likelihood and prior:

$$f_{\text{posterior}}(\theta|y) = \frac{f(y|\theta) \cdot f_{\text{prior}}(\theta)}{f(y)} \text{ where } f(y) = \int f(y|\theta) f_{\text{prior}}(\theta) d\theta$$

- We would like to generate the posterior distribution $f_{\text{posterior}}$
- We have likelihood $f(y|\theta)$ and an assumption for prior $f_{\text{prior}}(\theta)$
- For Metropolis-Hastings, we do not need the denominator $f(y)$; it cancels out in the MH ratio (see algorithm):

$$R(x^{(t)}, x^*) = \frac{f_{\text{posterior}}(x^*) g(x^{(t)}|x^*)}{f_{\text{posterior}}(x^{(t)}) g(x^*|x^{(t)})}$$


Metropolis algorithm - Example 2

(compare Givens and Hoeting, ex. 5.3)


- Consider Bayesian estimation of μ based on $N(\mu, 3^2/7)$ likelihood for μ and $\text{Cauchy}(5, 2)$ prior; observed mean=5.38
- The posterior density is proportional to product of likelihood and prior
- **Use MCMC to generate random samples following the posterior density**
- Based on these random samples, one can e.g.
 - determine posterior probability that $2 \leq \mu \leq 8$
 - determine mean and variance of posterior

Metropolis algorithm - Example 2

- We use starting value $x^{(0)}=0$, $s=1000$ iterations and following proposal distributions $g(\cdot|x^{(t)})$:
 $x^{(t)} + \text{Unif}[-0.2, 0.2]$, $x^{(t)} + \text{Unif}[-1, 1]$, $x^{(t)} + \text{Unif}[-8, 8]$
- **Sample path plots** show simulated values $x^{(t)}$ vs. iteration number t





Metropolis algorithm - Example 2



- Count “acceptance rate” (=proportion accepted proposals)
- Here: 98% 78% 18%
- Best results for 44% (uni-dim. case) to 23.4% (high dim. case) acceptance probability (theory based on normal target and proposal functions, see Givens and Hoeting, Chapter 7.3, for references about that)
- For multimodal functions lower acceptance probabilities might be good


Metropolis algorithm - Example 2

- Based on sample path plots, we might choose $\text{Unif}[-1,1]$ as proposal distribution
- Often, one wants to discard initial samples (**burn-in** period) which highly depend on starting value, e.g. **50 values + $x^{(0)}$**



Metropolis algorithm - Example 2


- For $s=10\ 000$ iterations and burn-in of 50, we obtain




- Monte Carlo estimate for $P(2 \leq \mu \leq 8)$ is 0.9967
(Monte Carlo standard error = $\sqrt{0.9967 * 0.0033 / 9950} = 0.0006$)
- Estimated mean = 5.26, standarddeviation = 0.99

Metropolis algorithm - Example 2

- Were $s=10\ 000$ iterations enough to ensure convergence?
- Can depend on the purpose ...
- E.g. for estimating $P(2 \leq \mu \leq 8)$
- One can monitor cusum/convergence plots showing estimate versus iterations (see Givens and Hoeting, ch.7.3.1.1)
- After 10 000 iterations



After 100 000 iterations



- After 10 000 iterations, we might not be happy with the left graph; we run longer and are happy with 100 000

Metropolis-Hastings with independent proposals

- Other proposal distributions g possible (not necessarily symmetric), e.g. independent proposals
- Proposal distribution depends not on previous value, $g(\cdot|x^{(t)})= g(\cdot)$
- The MH ratio is $R(x^{(t)}, x^*) = \frac{f(x^*) g(x^{(t)}|x^*)}{f(x^{(t)}) g(x^*|x^{(t)})} = \frac{f(x^*)/g(x^*)}{f(x^{(t)})/g(x^{(t)})}$
- A possible application is for Bayesian analysis (f is the posterior) with proposal distribution g being the prior distribution
- f/g is then the likelihood

Gibbs sampling

- Situation:
 - We want to sample a multivariate distribution $f(x_1, \dots, x_d)$ and this density is difficult to sample from
 - The conditional distributions for each single dimension i given fixed values $x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_d$ are easy to sample from
- Gibbs sampling generates a Markov chain converging to distribution f in this situation
- We sample one dimension in turn

Gibbs sampling

- The algorithm:

- o. A starting value $\mathbf{x}^{(0)} = (x_1^{(0)}, x_2^{(0)}, \dots, x_d^{(0)})$ is chosen

- 1. Generate $x_1^{(t+1)}$ following $f(x_1 | x_2^{(t)}, \dots, x_d^{(t)})$

- 2. Generate $x_2^{(t+1)}$ following $f(x_2 | x_1^{(t+1)}, x_3^{(t)}, \dots, x_d^{(t)})$

- ...

- i. Generate $x_i^{(t+1)}$ following $f(x_i | x_1^{(t+1)}, \dots, x_{i-1}^{(t+1)}, x_{i+1}^{(t)}, \dots, x_d^{(t)})$

- ...

- d. Generate $x_d^{(t+1)}$ following $f(x_d | x_1^{(t+1)}, \dots, x_{d-1}^{(t+1)})$

- Go back to 1. if more points needed

Gibbs sampling

- Example:

Let $f(x_1, x_2) = c \cdot \mathbf{1}\{x_1^2 + 1.8 \cdot x_1 x_2 + x_2^2 < 1\}$
be the uniform distribution on the ellipse

$$x_1^2 + 1.8 \cdot x_1 x_2 + x_2^2 < 1$$

- The conditional distribution for x_2 given x_1 is
a uniform distribution on the interval

$$(-0.9x_1 - \sqrt{1 - 0.19x_1^2}, -0.9x_1 + \sqrt{1 - 0.19x_1^2})$$

provided that term below root is positive

You can obtain these boundaries by solving $x_1^2 + 1.8 \cdot x_1 x_2 + x_2^2 - 1 = 0$ for x_2

- x_1 given x_2 has a similar distribution with x_2
instead of x_1 in the boundaries

