

# EM Algorithm, Stochastic Optimization

732A90

Computational Statistics

Krzysztof Bartoszek  
([krzysztof.bartoszek@liu.se](mailto:krzysztof.bartoszek@liu.se))

5 XII 2023 (U14)

Department of Computer and Information Science  
Linköping University

# Stochastic and combinatorial optimization

- So far: Unconstrained optimization
  - Predictor variables are continuous
  - Response function is differentiable
- We discussed Steepest descent, Newton, BFGS, CG
- But: predictors can be discrete  
(scheduling problems, travelling salesman)
- But: outcome can be discrete, noisy or multi-modal

Given a (large) set of states  $S$ , find

$$\min_{s \in S} f(s)$$

- Exhaustive search (shortest path algorithm)
- Often exhaustive search is NP-hard (TSP)
- Alternative: stochastic methods  
random search

Motivation from physics: cooling of metal

- Parameters:  
Energy of metal  
(decreasing, but not strictly monotonic)  
Temperature (decreasing)
- Aim: find global minimum energy

# Simulated annealing

```
1: Initialize  $k = 1$ ,  $\theta_0$ ,  $T(1) = T_0$ ,  $i = 0$ 
2: while  $k < k_{max} + 1$  do
3:   generate new state  $\xi$ , compute  $\Delta f = f(\xi) - f(\theta)$ 
4:   if  $\Delta f < 0$  then
5:     accept  $\xi$  ( $\theta_k = \xi$ )
6:   else
7:     accept  $\xi$  ( $\theta_k = \xi$ ) with  $P(\Delta f, T(k))$ 
8:   end if
9:    $k = k + 1$ 
   {WE CAN ALSO HAVE A SEARCH LOOP FOR EACH
    $T(k)$  level}
10: end while
```

N. Metropolis, A. Rosenbluth, M. Rosenbluth, A. Teller, E. Teller,  
Equations of state calculations by fast computing machines. J. Chem.  
Phys. 21:1087-1092, 1953.

# Simulated annealing

- [https://www.youtube.com/watch?v=iaq\\_Fpr4KZc](https://www.youtube.com/watch?v=iaq_Fpr4KZc)
- Generating new state:
  - Continuous: choose a new point a (random) distance from the current one
  - Discrete: similar or some rearrangement
- Selection probability: e.g  $\exp(-\Delta f(x)/T)$ : decreasing with  $f(x)$ , increasing with  $T$
- Temperature function: constant, proportional to  $k$ , or

$$T(k+1) = b(k)T(k), \quad b(k) = (\log(k))^{-1}$$

**Remember:** A smaller value is better than one on the path to the global minimum! Always keep track of smallest found.

# Simulated annealing: TSP example

Assume constant temperature

- 1: Choose initial configuration  $(Town_1, \dots, Town_n)$
- 2:  $k = 1$
- 3: **while**  $k < k_{max} + 1$  **do**
- 4:   Generate new configuration by rearrangement,  
           $(1, 2, 3, 4, 5, 6, 7, 8, 9) \rightarrow (1, 6, 5, 4, 3, 2, 7, 8, 9)$   
           $(1, 2, 3, 4, 5, 6, 7, 8, 9) \rightarrow (1, 7, 8, 2, 3, 4, 5, 6, 9)$
- 5:   Measure difference in path length  $(\Delta f)$  between old and new configuration
- 6:   **if** shorter path found **then**
- 7:     accept it
- 8:   **else**
- 9:     accept it with probability  $P(\Delta f)$
- 10:   **end if**
- 11:    $k++$
- 12: **end while**

# Genetic algorithm

- Inspiration from evolutionary theory: survival of the fittest
- Variables=genotypes
- Observation=organism, characterized by genetic code
- State space=population of organisms
- Objective function=fitness of organism

New points are obtained from old points by crossover and mutation, the population only retains the fittest organisms (with better objective function).

[https://en.wikipedia.org/wiki/List\\_of\\_genetic\\_algorithm\\_applications](https://en.wikipedia.org/wiki/List_of_genetic_algorithm_applications)



# Genetic algorithm

## Encoding points

- 1 Enumerate each element of the state space,  $S$
- 2 Code for observation  $i$  is binary representation of  $i$  (or something else)

## Mutation and recombination rules

*Crossover*: (1010 1110, 1100 0110)  $\longrightarrow$  1010 0110

*Inversion*: 11001011  $\longrightarrow$  11010011

*Mutation*: 11010111  $\longrightarrow$  1101**1**111

*Clone*: 11010111  $\longrightarrow$  11010111

# Genetic algorithm

- 1: Initialize  $i = 0$ , population  $\mathbf{P}_0$  and calculate fitness
- 2: **while** end criteria not met **do**
- 3:   Choose individuals  $\mathbf{T}_i$  from  $\mathbf{P}_i$  for reproduction
- 4:    $\mathbf{T}'_i =$  crossover between individuals from  $\mathbf{T}_i$
- 5:    $\mathbf{O}_i =$  randomly invert, mutate, clone  $\mathbf{T}'_i$
- 6:   Calculate fitness and obtain  $\mathbf{P}_i$  from  $\mathbf{O}_i$
- 7:    $i = i + 1$
- 8: **end while**

**ALWAYS KEEP BEST INDIVIDUAL**

# Genetic algorithm: TSP example

Encoding and crossover

- Encode tours as  $A_1, \dots, A_n$  but

Parent 1: FAB|ECGD      Parent 2: DEA|CGBF  
Child: FAB|CGBF      Child: DEA|ECGD

Instead

- 1 Remove FAB from DEACGBF  $\rightarrow$  DECG.  
Child becomes FABDECG.
- 2 Second child will be by taking prefix from Parent 2:  
DEAFBCG

# Genetic algorithm: Mutations


- If a population is small and only crossover: the input domain becomes limited and may converge to a local minimum.
- Large initial populations are computationally heavy.
- Mutations allow one to explore more of  $S$ : jump out of local minimum.
- In TSP: mutation move a city in the tour to another position.
- Reproduction: Among  $m$  tours selected at step 2, two best are selected for reproduction, two worst replaced by children.
- If  $m$  is large, some tours might never be parents, global solution may be missed. Random chance of reproduction?
- Mutation probability is usually small (unless you want to jump wildly)

# Genetic algorithm: example

Minimize

$$f(x) := \frac{x^2}{e^x} - 2 \exp(-(9 \sin x)/(x^2 + x + 1))$$

- `fitness(x) = -f(x)`,
- `crossover(x,y) = (x + y)/2`,
- `mutate(x) = x +  $\epsilon$ ; =  $x^2 \bmod 30$`   
( `mod`  $\equiv$  `%%`—integer division )
- smallest fitness individual removed



- mutation probability: 0.5 worked well; 0.1: cannot change;  
0.9: cannot stabilize

# EM algorithm

**Fundamental algorithm** of computational statistics!

Model depends on the data which are observed (known) **Y** and **latent** (unobserved) data **Z**.

The data's (**both Y's and Z's**) distribution depends on some parameters  $\theta$ .

**AIM:** Find MLE of  $\theta$ .

- All data is known: Apply unconstrained optimization (discussed in Lecture 2)
- Unobserved data
  - **Sometimes** it is possible to look at the marginal distribution of the observed data.
  - Otherwise: **EM algorithm**

# EM algorithm

Let

$$Q(\theta, \theta^k) = \int \log p(\mathbf{Y}, \mathbf{z}|\theta) p(\mathbf{z}|\mathbf{Y}, \theta^k) d\mathbf{z} = \mathbb{E} \left[ \log \text{lik}(\theta|\mathbf{Y}, \mathbf{Z}) | \theta^k, \mathbf{Y} \right]$$

- 1:  $k = 0, \theta^0 = \theta^0$
- 2: **while** Convergence not attained **and**  $k < k_{max} + 1$  **do**
- 3:   **E-step:** Derive  $Q(\theta, \theta^k)$
- 4:   **M-step:**  $\theta^{k+1} = \operatorname{argmax}_{\theta} Q(\theta, \theta^k)$
- 5:    $k++$
- 6: **end while**

**Example:** Normal data with missing values (but here analytical approach is also possible)

732A90\_ComputationalStatisticsHT2022\_Lecture06codeSlide16.R



## EM algorithm: R

```
> Y<-rnorm(100)
> Y[sample(1:length(Y),20,replace=FALSE)]<-NA
> EM.Norm(Y,0.0001,100)
[1] 1.0000 0.1000 -997.5705
[1] 0.1341894 1.3227095 -128.2789837
[1] -0.03897274 1.38734070 -126.86036252
[1] -0.07360517 1.39307050 -126.80801589
[1] -0.08053165 1.39392861 -126.80593837
[1] -0.08191695 1.39408871 -126.80585537
> mean(Y,na.rm=TRUE)
[1] -0.08226328
> var(Y,na.rm=TRUE)
[1] 1.411775
```

Notice: can be done by studying marginal distribution of observed data.

# EM algorithm: Applications

**Mixture models**  $Z$  is a latent variable,  $P(Z = k) = \pi_k$

- Mixed data comes from different sources (e.g. for regression, classification)
- Clustering
  - 1 Density in each cluster is normally distributed.
  - 2 Cluster label is latent (we do not know what are the chances an observation is from the given cluster)

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x | \vec{\mu}_k, \Sigma_k) \quad (\text{informally})$$

Direct MLE leads to numerical problems.

Introduce latent class variables and use EM.

# EM algorithm: Gaussian mixtures

- 1: Initialize  $r = 0$ ,  $(\vec{\mu}_1^0, \dots, \vec{\mu}_K^0)$ ,  $(\Sigma_1^0, \dots, \Sigma_K^0)$  **input:** data  $\vec{x}_1, \dots, \vec{x}_n$ ;  $K$  clusters  
2: **while** end criteria not met **do**  
3:   **E-step** ( $\bar{z}_{jk}$ : *responsibility* of cluster  $k$  explaining  $\vec{x}_j$ )

$$\bar{z}_{jk} = \frac{\mathcal{N}(\vec{x}_j | \vec{\mu}_k^r, \Sigma_k^r)}{\sum_{i=1}^K \pi_i \mathcal{N}(\vec{x}_j | \vec{\mu}_i^r, \Sigma_i^r)}, \quad j = 1, \dots, n$$

- 4:   **M-step** for  $k = 1, \dots, K$

$$n_k = \sum_{j=1}^n \bar{z}_{jk}, \quad \pi_k^r = n_k/n,$$

$$\vec{\mu}_k^r = (n_k)^{-1} \sum_{j=1}^n \bar{z}_{jk} \vec{x}_j, \quad \Sigma_k^r = (n_k)^{-1} \sum_{j=1}^n \bar{z}_{jk} (\vec{x}_j - \vec{\mu}_k^r)(\vec{x}_j - \vec{\mu}_k^r)^T$$

- 5:   evaluate log-likelihood

$$\log \mathcal{L}(\{\pi_k^r, \vec{\mu}_k^r, \Sigma_k^r\}_{k=1}^K) = \sum_{j=1}^n \log \left( \sum_{k=1}^K \pi_k^r \mathcal{N}(\vec{x}_j | \vec{\mu}_k^r, \Sigma_k^r) \right)$$


- 6:    $r = r + 1$

See also Ch. 7.10.7, *EM for Mixture of Gaussians*,

- 7: **end while**

O. Simeon, *Machine Learning for Engineers*, Cambridge University Press, 2023 (there supervised log-loss optimized)

## Gaussian mixtures: example



$$P(1) = P(2) = P(3) = P(4) = 0.25$$

- ➊ draw class  $Z \in \{1, 2, 3, 4\}$  uniformly
- ➋ draw normal distribution  $\mathcal{N}(Z, Z)$  with density  $\phi_{Z,Z}(\cdot)$

We can write the mixture density as

$$f(x) = 0.25\phi_{1,1}(x) + 0.25\phi_{2,2}(x) + 0.25\phi_{3,3}(x) + 0.25\phi_{4,4}(x).$$

# Stochastic gradient descent (SGD)—2016 slides

Goal : maximize  $g(\theta) = N^{-1} \sum_{j=1}^N g_j(\theta)$  for large  $N$

$\gamma_i$ : learning rate

1:  $\theta_0 = \theta_0, i = 0$

2: **while** stopping criterion not met **do**

3:   subset  $S_i \subseteq \{1, \dots, N\}$ , with  $|S_i| = S$

4:   update:

$$\theta_{i+1} = \theta_i + \frac{\gamma_i}{S} \sum_{j \in S_i} g_j(\theta)$$

5:    $i++$

6: **end while**

Ch. 5.11; O. Simeone, Machine Learning for Engineers, 2023,  
Cambridge University Press, Cambridge

# Variational inference (VI)

$x$ : observed variable;  $z$ : latent variable;

Model:  $p_\psi(x, z) = p_\psi(x|z)p_\psi(z)$

$\psi$ : unknown parameters

likelihood :  $L(\psi) = p_\psi(x) = \int p_\psi(x, z)dz$  not tractable

$q_\theta(\cdot)$ : known parametric family of distributions

$q_{\theta_\psi}(z) \approx p_\psi(z|x)$  in Kullback-Leibler divergence sense,

$$q_{\theta_\psi}^*(z) = \arg \min_{\theta_\psi} \text{KL}(q_{\theta_\psi}(z), p_\psi(z|x))$$

# Variational inference (VI)

$$\begin{aligned}\text{KL}(q_\theta(z), p_\psi(z|x)) &:= \int \log q_\theta(z) q_\theta(z) dz - \int \log p_\psi(z|x) q_\theta(z) dz \\ &= \int \log q_\theta(z) q_\theta(z) dz - \int \log p_\psi(z, x) q_\theta(z) dz + \log p_\psi(x) \\ &\geq \int \log q_\theta(z) q_\theta(z) dz - \int \log p_\psi(z, x) q_\theta(z) dz =: (-1) ELBO(q_\theta)\end{aligned}$$

*ELBO*: evidence lower bound, as  $\log p_\psi(x) \geq ELBO(q_\theta)$

estimate  $ELBO(q_\theta)$  using that  $q_\theta(z)$  is a distribution  
maximize  $ELBO(q_\theta)$  w.r.t.  $\theta$  (minimizes KL),

$$p_\psi(x) = p_\psi(z|x)p_\psi(z, x) \approx q_{\theta_\psi^*}(z)p_\psi(z, x)$$

Random walk over the state space in search of minimum

- ① Follow decreasing path
- ② **BUT** with a certain probability go to higher values, to avoid local minima traps.
- ③ **Never forget** best found conformation!
- ④ Simulated annealing, Genetic algorithm, **EM algorithm**, SGD, VI